

Class 07 : Machine Learning 1

Kyle Wittkop (A18592410)

Table of contents

Background	1
K-means, cluster	3
Hierarchical clustering	6
Principal component analysis (PCA)	9
PCA of UK food data	9
Heatmap	13
PCA to the rescue	14
Digging Deeper - Variable loadings	17


Background

Today we will begin our exploration of important machine learning methods, as a focus on **clustering** and **dimensionality reduction**.

To start testing these methods, lets make up some sample data dots to cluster where we know what the answer should be.

```
hist(rnorm(3000, mean=10))
```

Histogram of rnorm(3000, mean = 10)



Q. Can you generate 30 numbers centered around +3 taken atrandom drom a normal distribution

```
rnorm(30, mean=3)
```


```
[1] 3.573519 3.128351 3.555634 4.049601 1.830289 1.206520 3.492734 1.933054  
[9] 4.190879 2.823972 1.218854 3.594317 1.426486 1.563269 2.388148 3.786902  
[17] 2.569394 3.186117 4.055719 2.732841 1.534143 3.430038 3.064831 1.347411  
[25] 3.015511 2.325580 4.872557 3.404735 4.887299 3.779791
```

```
rnorm(30,mean=-3)
```

```
[1] -3.1029469 -3.6234436 -2.7447262 -2.2855722 -3.5020828 -3.0139025  
[7] -1.4996169 -3.1542440 -4.3546607 -4.6675254 -3.5769224 -3.0029604  
[13] -2.6449638 -2.5976411 -3.1652244 -2.6475485 -4.9079405 -3.8105808  
[19] -3.8134771 -3.2419254 -1.9414440 -1.0507329 -2.3860012 -3.7321387  
[25] -3.0312606 -3.1778409 -2.4186611 -1.1865320 -1.8811112 -0.9851372
```

```
tmp <- c(rnorm(30, mean=3),  
rnorm(30,mean=-3))
```

```
x <- cbind(x=tmp, y=rev(tmp))  
plot(x)
```



K-means, cluster

The main function in Base R, for K means clustering is called `Kmeans()`. Lets try it out.

```
k <- kmeans(x, centers = 2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

          x           y
1  3.258531 -3.012321
2 -3.012321  3.258531

```

Clustering vector:

```
Within cluster sum of squares by cluster:  
[1] 65.75109 65.75109  
(between_SS / total_SS =  90.0 %)
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. What component of your result object has the cluster centers

k\$centers

```

          x           y
1  3.258531 -3.012321
2 -3.012321  3.258531

```

Q. What component of your result object has the cluster size

k\$size


[1] 30 30

Q. What component of your result object has the cluster membership vector (i.e the main clustering result: which points are in which cluster)

k\$cluster

Q. Plot the result of cluster (i.e)

```
plot(x, col= k$cluster)
points(k$centers, col = "blue", pch=15, cex=2)
```



Q. Can you run k means again and cluster into 4 groups and plot the results

```
k <- kmeans(x, centers = 4)
k
```

K-means clustering with 4 clusters of sizes 13, 17, 17, 13

Cluster means:

	x	y
1	-2.093082	2.575785
2	-3.715269	3.780631
3	3.780631	-3.715269
4	2.575785	-2.093082

Clustering vector:

```
[1] 4 3 3 3 4 4 3 3 4 4 3 3 4 3 4 3 4 3 4 3 4 3 3 3 3 3 3 4 4 1 1 2 2 2 2 2 2
[39] 1 2 1 2 1 1 2 1 2 2 1 1 2 2 1 1 2 2 2 1
```

Within cluster sum of squares by cluster:


```
[1] 14.73038 20.94155 20.94155 14.73038
```

(between_SS / total_SS = 94.6 %)

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
plot(x, col= k$cluster)  
points(k$centers, col = "blue", pch=15, cex=2)
```



Key point K-means will always return clustering that we ask for, this is the k or centers in K means.

```
k$tot.withinss
```


```
[1] 71.34387
```

Hierarchical clustering

The main function to do this in base R is called hclust `hclust()`. One of the main differences with respect to the k means function, is that you cannot put your input data directly. it needs a distance matrix or a dissimilarity matrix. we can get this from lots of places including the `dist` function,

```
d <- dist(x)
hc <- hclust(d)
plot(hc)
```

Cluster Dendrogram




```
d
hclust (*, "complete")
```

We can cut the deprogram, at a given height to yield our clusters

```
plot(hc)
abline(h=10, col="red")
```

Cluster Dendrogram



```
d  
hclust (*, "complete")
```


```
grps<-cutree(hc, h=10)
```

```
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2  
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Q, Plot out data 'x' closed by clustering resulting from hclust (), and cutree

```
plot(x, col=grps)
```



Principal component analysis (PCA)

PCA is a popular dimensionality reduction technique that is widely used in bioinformatics

PCA of UK food data

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this question?

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)  
dim(x)
```

```
[1] 17 5
```

It looks like the row names were not set properly. We can fix this.

```

rownames(x) <- x[,1]
x <- x[,-1]
x

```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

A Better way to do this is to fix the row names assignment at import times :

```

x <-read.csv(url, row.names = 1)
x


```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674


Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q3: Changing what optional argument in the above barplot() function results in the following plot?

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```

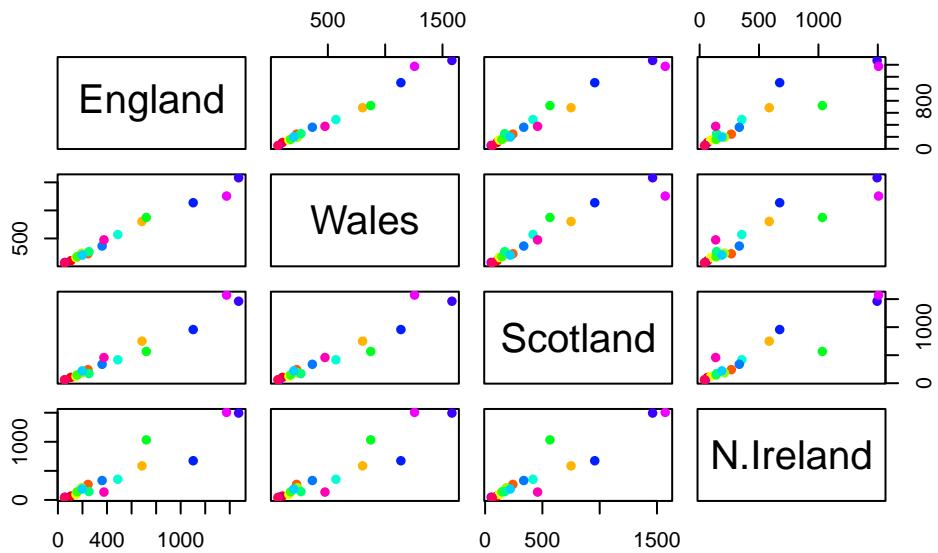


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: We can use the pairs() function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```

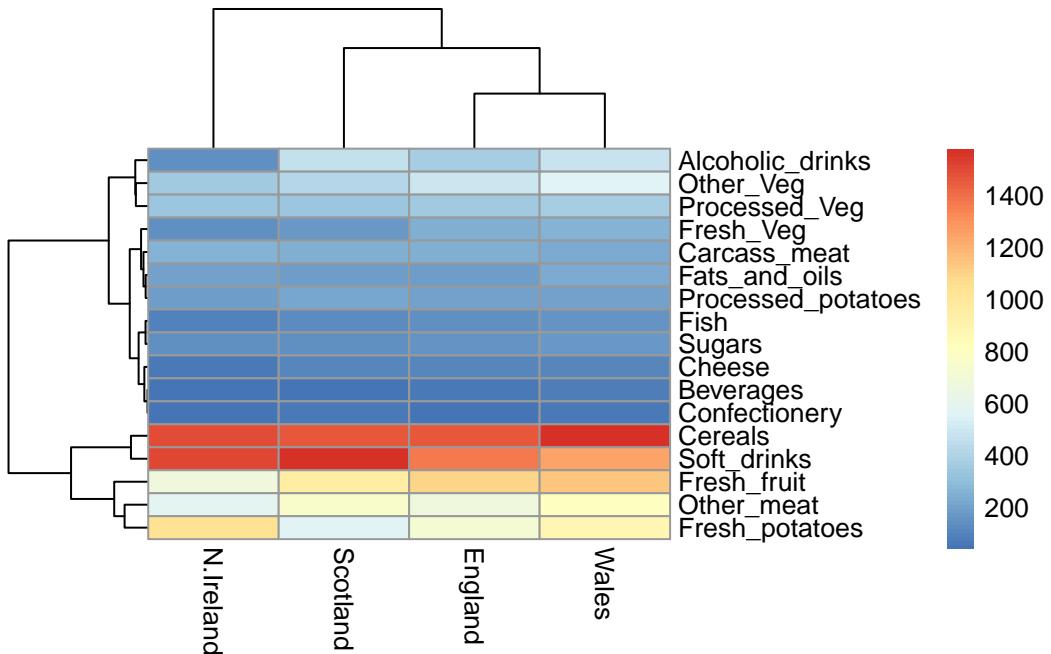


each point represents the data surrounding a food consumed by the 4 countries. If a given point lies on the diagonal for a given plot, it showed that the point is shared in value between the two countries plotted against each other. If it is off the diagonal, it means there is different values for these foods in these two countries.

Heatmap

we can install the ***pheatmap*** package with the `install.packages` command that we used the previously. Remember that we always run this in console and not a code chunk in our quarto document

```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Based on the pairs and heat map figures, it appears that England, Wales, and Scotland cluster together suggesting their food consumption patterns are similar. We can fairly easily tell the difference between Ireland and the other countries, however the pairs shows this difference more effectively as seen in the disorder of points along the diagonal when comparing to the other countries.

Of all these plots, really only the pairs plot was useful. This however took a bit of work to interpretive and will not scale when looking at much bigger data sets.

PCA to the rescue

The main function in “base R” for PCA is called `prcomp()`

```
pca <- prcomp ( t(x) )
summary(pca)
```

Importance of components:

PC1	PC2	PC3	PC4
-----	-----	-----	-----

```

Standard deviation      324.1502 212.7478 73.87622 2.7e-14
Proportion of Variance  0.6744    0.2905   0.03503 0.0e+00
Cumulative Proportion   0.6744    0.9650   1.00000 1.0e+00

```

Q. How much variance is captured in the first PC?

67.44% of variance is captured in the first PC

Q. How many PCs do I need to capture to achieve at least 90% of the total variance in the data set?

We need to capture 2 PCs in order to achieve a 96.5% of variance shown in cumulative proportions.

Q Plot out main PCA results Folks can call this different things depending on their field of study for example PC plot.

```
attributes(pca)
```

```

$names
[1] "sdev"      "rotation"  "center"    "scale"     "x"

$class
[1] "prcomp"

```

To generate our plot we want pca\$x


```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	1.612425e-14
Wales	-240.52915	-224.646925	-56.475555	4.751043e-13
Scotland	-91.86934	286.081786	-44.415495	-6.044349e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.145386e-13


```

my_cols <- c("orange","red","blue","green")
plot(pca$x [,1], pca$x [,1])

```



```
library(ggplot2)
ggplot(pca$x) +
  aes(x = PC1, y = PC2, label = rownames(pca$x)) +
  geom_point(size = 3, col=my_cols) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()
```



Digging Deeper - Variable loadings

How do the original variables (17 different foods) contribute to our new PCS?

```
ggplot(pca$rotation) +
  aes(x = PC1,
      y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```

