

Class08

Kyle Wittkop (A18592410)

Background

In today's class we will apply the methods and techniques clustering and PCA to help make sense of real world breast cancer data

Data import

We start by importing our data. It is a CSV file so we will use the `read.csv()` function

```
#read.csv("WisconsinCancer.csv")
fna.data <- "WisconsinCancer.csv"

wisc.df <- read.csv(fna.data, row.names=1)
```

Make sure to remove the first `diagnosis` column - I don't want to use this for my machine learning models. We will use it later to compare our results to the expert diagnosis.

```
head(wisc.df, 4)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	

84348301	0.14250	0.28390	0.2414	0.10520		
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se	
842302	0.2419		0.07871	1.0950	0.9053	8.589
842517	0.1812		0.05667	0.5435	0.7339	3.398
84300903	0.2069		0.05999	0.7456	0.7869	4.585
84348301	0.2597		0.09744	0.4956	1.1560	3.445
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399		0.04904	0.05373	0.01587
842517	74.08	0.005225		0.01308	0.01860	0.01340
84300903	94.03	0.006150		0.04006	0.03832	0.02058
84348301	27.23	0.009110		0.07458	0.05661	0.01867
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38	17.33	
842517	0.01389		0.003532	24.99	23.41	
84300903	0.02250		0.004571	23.57	25.53	
84348301	0.05963		0.009208	14.91	26.50	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302		184.60	2019.0	0.1622	0.6656	
842517		158.80	1956.0	0.1238	0.1866	
84300903		152.50	1709.0	0.1444	0.4245	
84348301		98.87	567.7	0.2098	0.8663	
	concavity_worst	concave.points_worst	symmetry_worst			
842302		0.7119	0.2654	0.4601		
842517		0.2416	0.1860	0.2750		
84300903		0.4504	0.2430	0.3613		
84348301		0.6869	0.2575	0.6638		
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				

```
wisc.data <- wisc.df[,-1]

diagnosis <- wisc.df$diagnosis

nrow(wisc.data)
```

[1] 569

```
dim(wisc.data)
```

```
[1] 569 30
```

```
table(diagnosis)
```

```
diagnosis
```

```
  B   M
```

```
357 212
```

```
coln <- colnames(wisc.data)
```

```
length(grep("_mean",coln))
```

```
[1] 10
```

Q1. How many observations are in this dataset?

569

Q2. How many of the observations have a malignant diagnosis?

212

Q3. How many variables/features in the data are suffixed with _mean?

10

Principal component analysis

The main function is `prcomp()` and we want to make sure we set the optional argument `Scale=True`

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean

	8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se	
	6.279761e-02	4.051721e-01	1.216853e+00
	perimeter_se	area_se	smoothness_se
	2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se	
	2.547814e-02	3.189372e-02	1.179614e-02
	symmetry_se	fractal_dimension_se	radius_worst
	2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst	
	2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst	
	1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst	
	1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

```
wisc.pr <- prcomp(wisc.data, scale=TRUE, )
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

44.27%

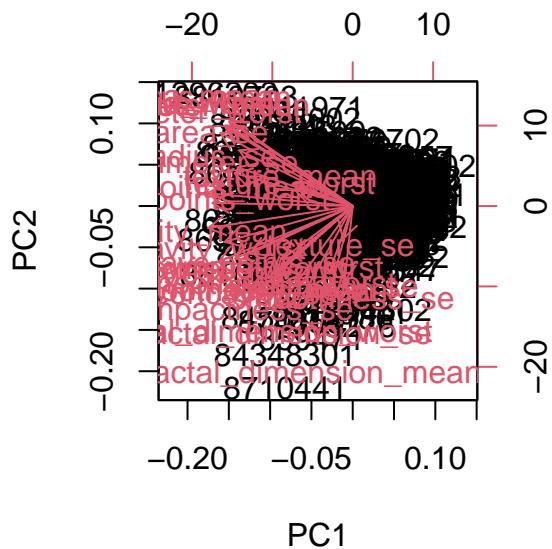
Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

PC3

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

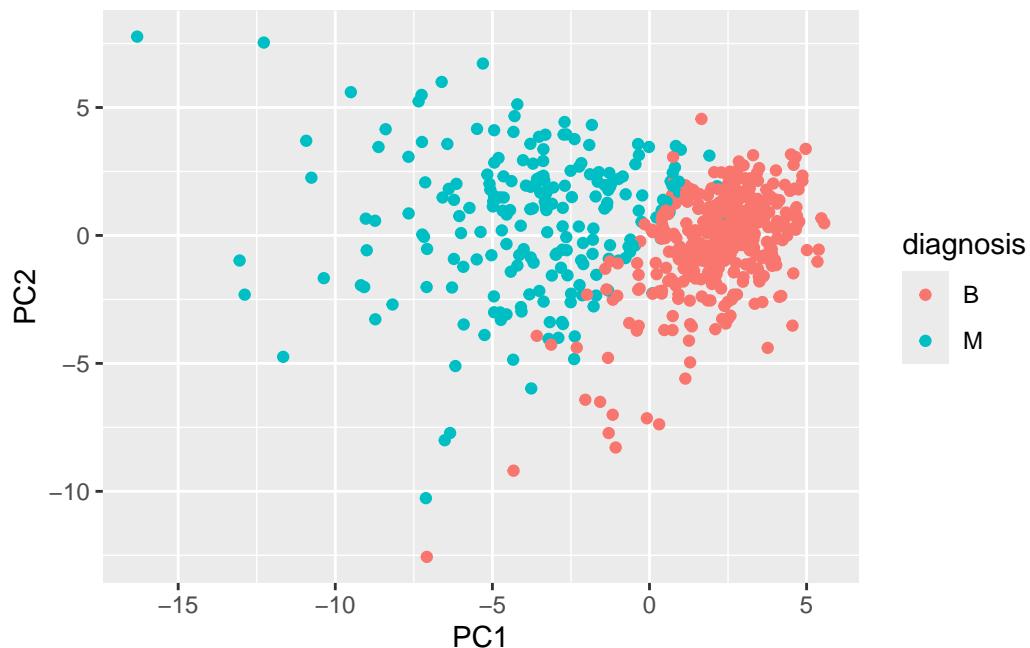
PC7

```
biplot(wisc.pr)
```

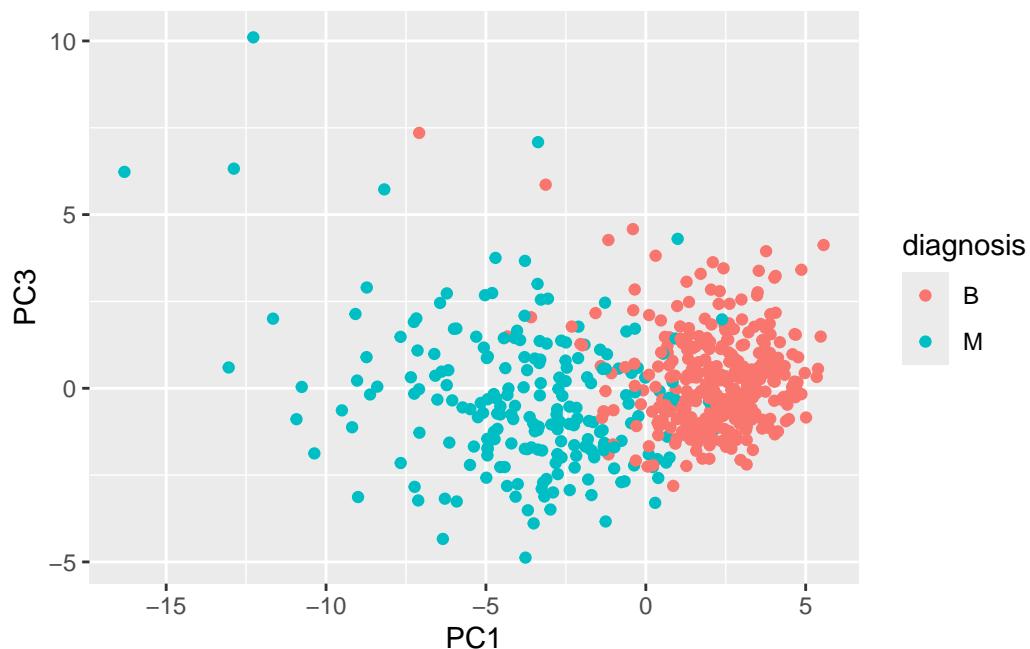


```
library(ggplot2)

ggplot(wisc.pr$x) +
  aes(PC1,PC2, col=diagnosis) +
  geom_point()
```



```
ggplot(wisc.pr$x) +
  aes(PC1,PC3, col=diagnosis) +
  geom_point()
```



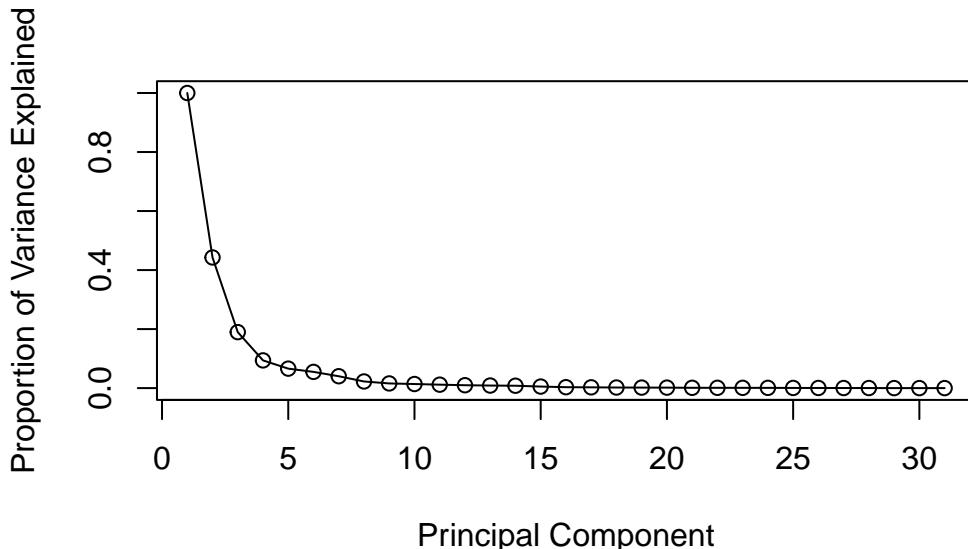
```
pr.var <- wisc.pr$sdev^2  
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

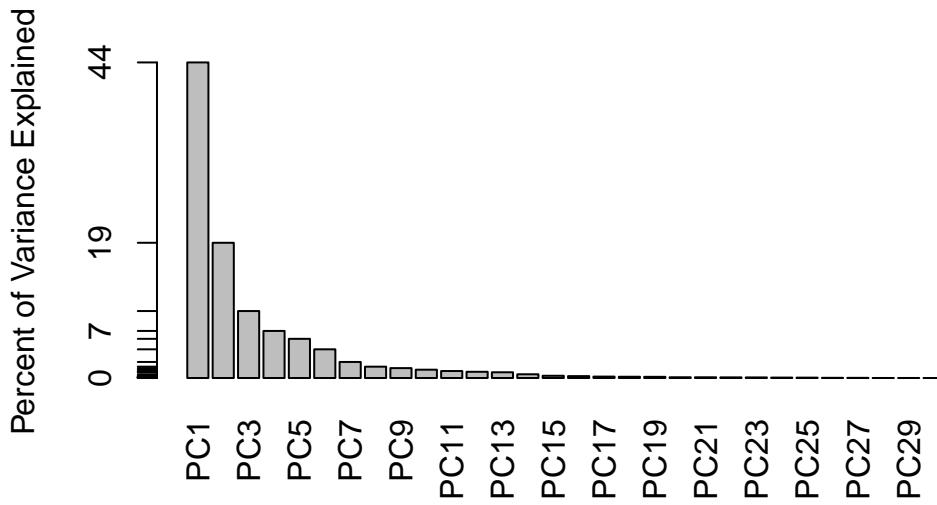
```
total.var <- sum(wisc.pr$sdev^2)
```

```
pve <- pr.var / total.var
```

```
plot(c(1,pve), xlab = "Principal Component",  
     ylab = "Proportion of Variance Explained",  
     ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Percent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
wisc.pr$rotation["concave.points_mean", 1]
```

```
[1] -0.2608538
```

```
data.scaled <- scale(wisc.data)

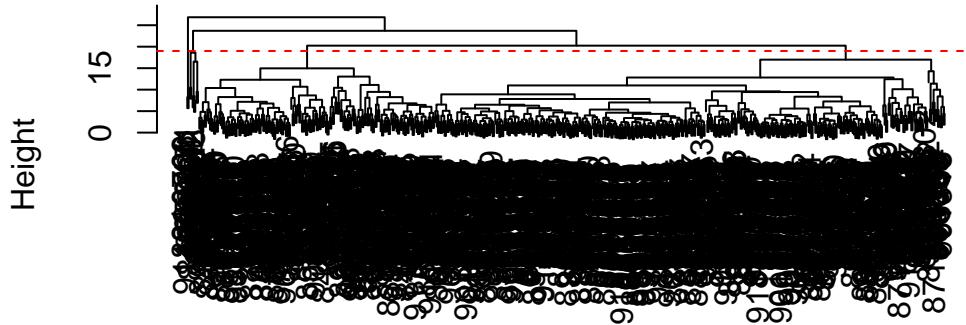
data.dist <- dist(data.scaled)

wisc.hclust <- hclust(data.dist, "complete")
```

Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
```

```
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

the height is 19 to give us 4 clusters.

You can also use the argument `k=4` in `cutree()` function instead of `h`.

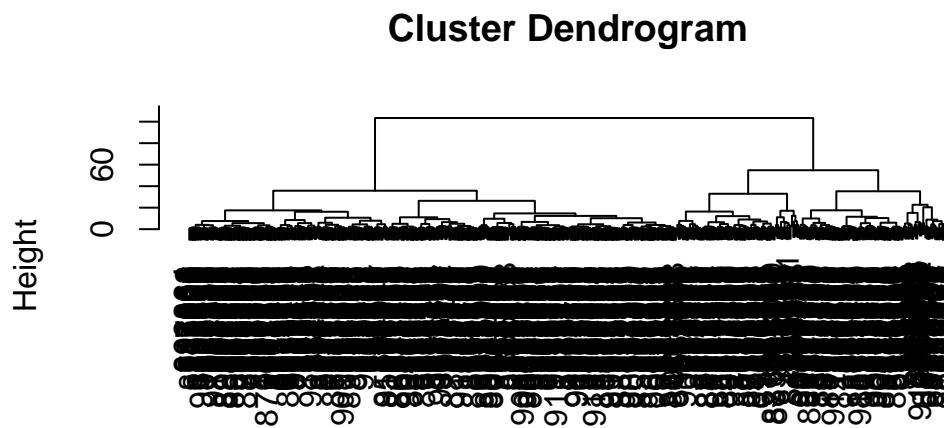
Combining Methods

Here we will take our PCA results and use those as input for clustering. in other words our `wisc.pr$x` scores that we plotted above (the main output from PCA - How the data lie in our new principal component axis/variables) and use a subset of these PCs that capture the most variance as input for `hclust()`

```

pc.dist <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust<-hclust(pc.dist,"ward.D2")
plot(wisc.pr.hclust)

```



```

pc.dist
hclust (*, "ward.D2")

```

I want to know how clustering into groups with values of 1 or 2 coorespond to the diagnosis

```

wisc.pr.hclust<-hclust(pc.dist,"ward.D2")

grps <- cutree(wisc.pr.hclust, k=2)
table(grps)

```

```

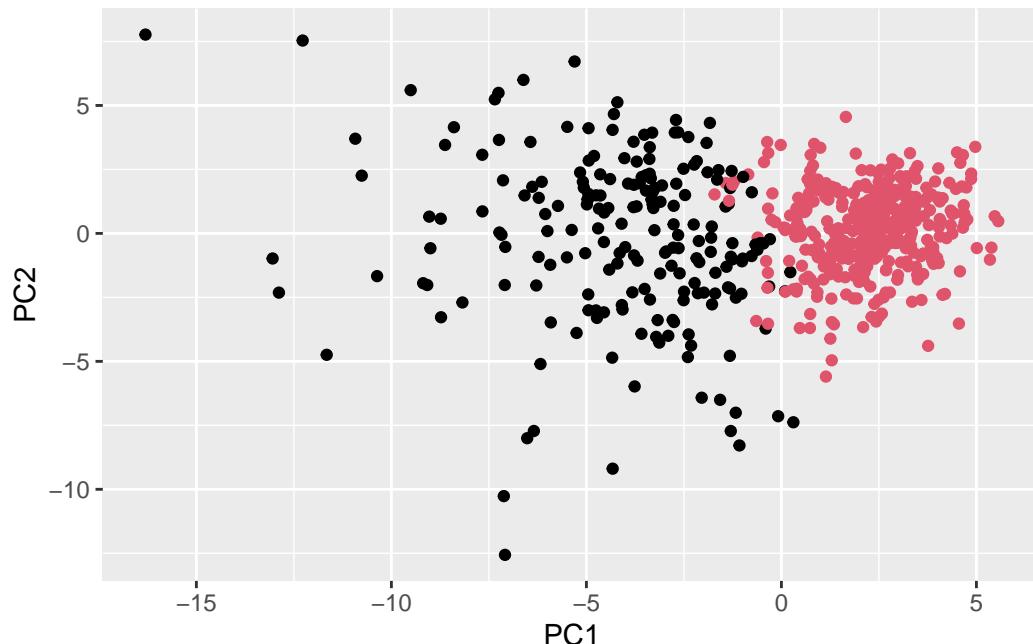
grps
 1 2
203 366

```

```
table(grps, diagnosis)
```

grps	B	M
1	24	179
2	333	33

```
ggplot(wisc.pr$x) +
  aes(PC1, PC2) +
  geom_point(col=grps)
```



My clustering **groups 1** of M diagnosis (179) and my clustering **group 2** are mostly “B” diagnosis

24 False positives 179 true positives

Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

the “ward.D2” method

Q13. How well does the newly created hclust model with two clusters separate out the two “M” and “B” diagnoses?

```
dist7<-dist(wisc.pr$x[,1:7])
wisc.pr.hclust <- hclust(dist7, method="ward.D2")

wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)

table(wisc.pr.hclust.clusters, diagnosis)
```

```

diagnosis
wisc.pr.hclust.clusters   B   M
                           1 28 188
                           2 329 24

```

the new hclust model with two clusters appear to separate them out similar to the previous one despite considering up to PC7

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters)
```

```
wisc.hclust.clusters
 1   2   3   4
177  7 383  2
```

```
table(wisc.pr.hclust.clusters)
```

```
wisc.pr.hclust.clusters
 1   2
216 353
```

Sensitivity/Specificity - Prediction

```

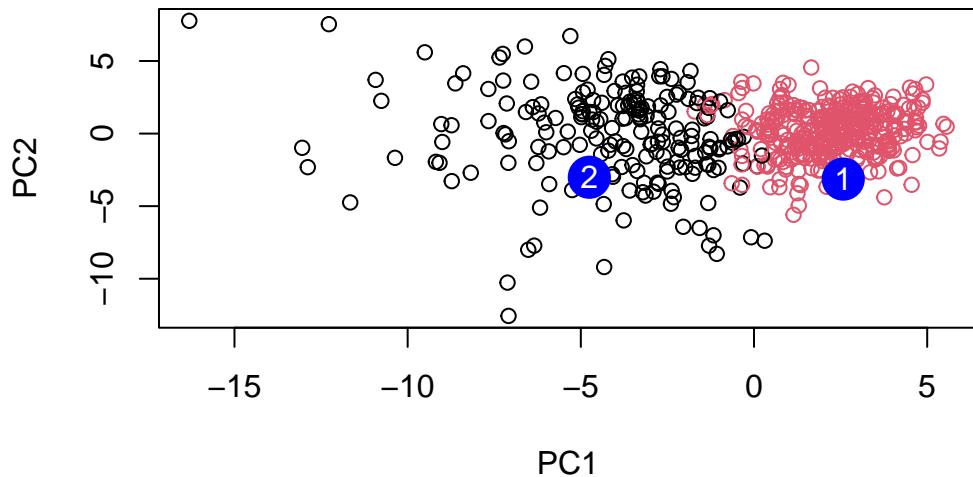
new <- read.csv("new_samples.csv")
npc <- predict(wisc.pr, newdata=new)
npc

```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	

```
[2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
      PC21      PC22      PC23      PC24      PC25      PC26
[1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
      PC27      PC28      PC29      PC30
[1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
[2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results?

Patients cluster under 2 we should prioritize for follow up based on results.