

Class09 - Candy Project

Kyle Wittkop (A18592410)

Table of contents

Importing candy data	1
Exploratory analysis	5
Overall Candy Rankings	7
6 Exploring the coorelation structure	14
Principal Component Analysis	15
Summary	20
Optional extension questions	21

Importing candy data

```
candy_file <- "candy-data.csv"

candy = read.csv(candy_file, row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crisped	rice	wafer
100 Grand	1	0	1			0	0		1
3 Musketeers	1	0	0			0	1		0
One dime	0	0	0			0	0		0
One quarter	0	0	0			0	0		0
Air Heads	0	1	0			0	0		0
Almond Joy	1	0	0			1	0		0
	hard	bar	pluribus	sugar	percent	price	percent	win	percent
100 Grand	0	1	0	0.732		0.860	66.97173		
3 Musketeers	0	1	0	0.604		0.511	67.60294		
One dime	0	0	0	0.011		0.116	32.26109		
One quarter	0	0	0	0.011		0.511	46.11650		

```
Air Heads      0   0       0       0.906      0.511    52.34146
Almond Joy     0   1       0       0.465      0.767    50.34755
```

```
nrow(candy)
```

```
[1] 85
```

```
sum(candy$fruity)
```

```
[1] 38
```

Q1. How many different candy types are in this dataset?

```
85
```

Q2. How many fruity candy types are in the dataset?

```
38
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

```
filter, lag
```

The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```

```
candy %>%
  filter(row.names(candy)=="Twix") %>%
  select(winpercent)
```

```
winpercent
Twix    81.64291
```

```
candy["Twix", ]$winpercent
```

```
[1] 81.64291
```

Q3. What is your favorite candy (other than Twix) in the dataset and what is its winpercent value?

```
candy["Air Heads", ]$winpercent
```

```
[1] 52.34146
```

Air Head win percent = 52.34146

Q4. What is the winpercent value for “Kit Kat”?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

76.7686

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

49.6535

```
library("skimr")
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12

Column type frequency:	
numeric	12
<hr/>	
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

Yes, Win percent seems to be on a different scale, all the other values are between 0 and 1 but winpercent is much larger.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

candy\$chocolate

```
[1] 1 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 0 1 1 1 0 0 0 1 1 0 1 1 1  
[39] 1 1 1 0 1 1 0 0 0 1 0 0 0 1 1 1 1 1 0 1 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 0 1 1  
[77] 1 1 0 1 0 0 0 0 1
```

```
table(candy$chocolate)
```

0 1
48 37

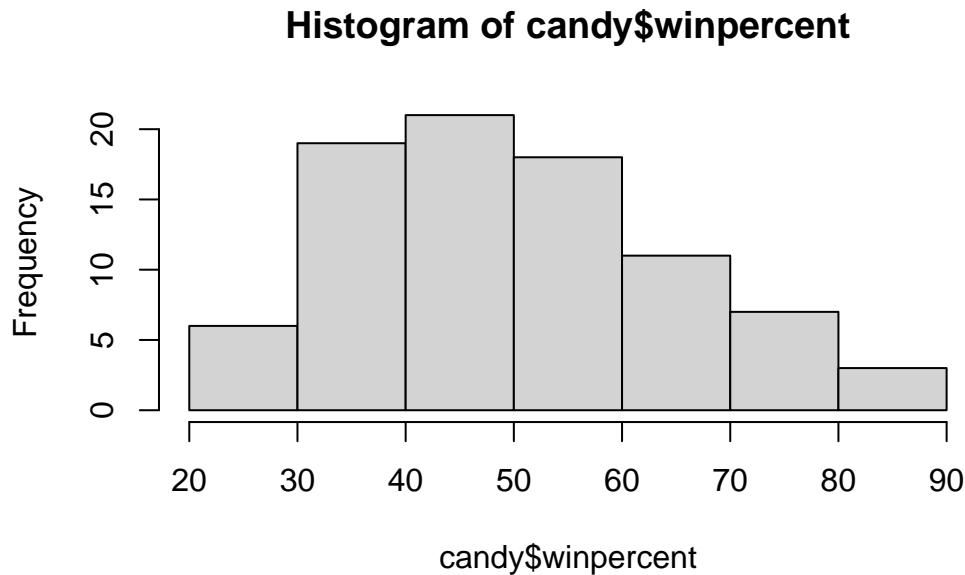
I think 0 means the candy doesn't contain chocolate and 1 the candy does contain chocolate.

Exploratory analysis

Q8. Plot a histogram of winpercent values

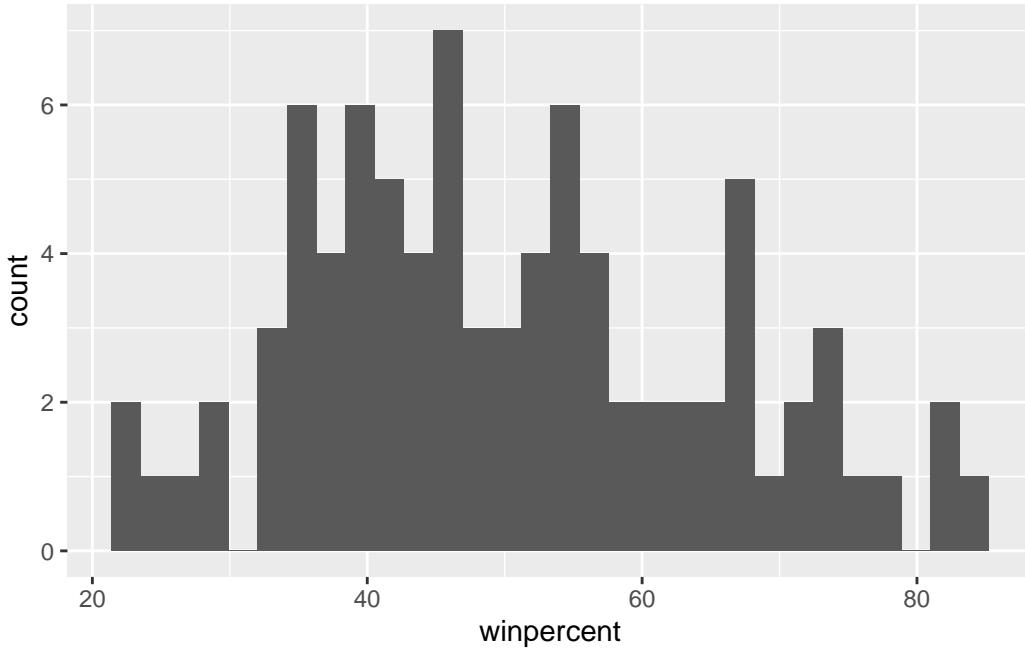
Base R:

```
hist(candy$winpercent)
```



ggplot:

```
library(ggplot2)
ggplot(candy,aes(winpercent)) + geom_histogram(bins=30)
```



Q9. Is the distribution of winpercent values symmetrical?

No, the distribution seems to be right skewed.

Q10. Is the center of the distribution above or below 50%?

```
median(candy$winpercent)
```

[1] 47.82975

median of winpercent = 47.82975 seems to be below 50%

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
chocwin<-candy$winpercent [as.logical(candy$chocolate)]
fruitwin<-candy$winpercent [as.logical(candy$fruity)]

mean(chocwin)
```

[1] 60.92153

```
mean(fruitwin)
```

```
[1] 44.11974
```

mean of chocolate win percent = 60.92153 mean of fruity win percent = 44.11974
on average chocolate is higher ranked than fruity.

Q12. Is this difference statistically significant?

```
t.test(chocwin,fruitwin)
```

Welch Two Sample t-test

```
data: chocwin and fruitwin
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

p value = 2.871e-08 which is less than 0.05 so it is considered significant.

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
candy %>% arrange(winpercent) %>% head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	
Nik L Nip	0	1	0	0	0	0	
Boston Baked Beans	0	0	0	1	0	0	
Chiclets	0	1	0	0	0	0	
Super Bubble	0	1	0	0	0	0	
Jawbusters	0	1	0	0	0	0	
	crispedrice	wafer	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0	1	0.197	0.976	

Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
winpercent						
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

the 5 least liked candy types are 1) Nik L Nip 2) Boston Bakes beans 3) Chiclets 4) Super Bubble 5) Jawbusters

Q14. What are the top 5 all time favorite candy types out of this set?

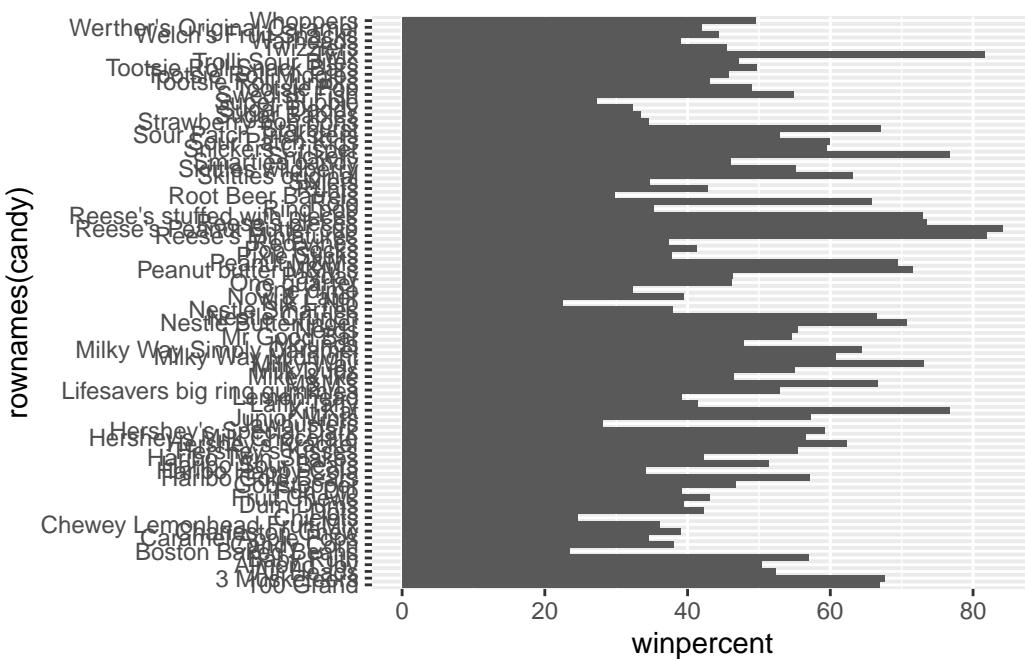
```
candy %>% arrange(desc(winpercent)) %>% head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Peanut Butter cup	1	0	0	1	0	
Reese's Miniatures	1	0	0	1	0	
Twix	1	0	1	0	0	
Kit Kat	1	0	0	0	0	
Snickers	1	0	1	1	1	
	crispedrice	wafer	hard	bar	pluribus	sugarpercent
Reese's Peanut Butter cup	0	0	0	0	0	0.720
Reese's Miniatures	0	0	0	0	0	0.034
Twix	1	0	1	0	0	0.546
Kit Kat	1	0	1	0	0	0.313
Snickers	0	0	1	0	0	0.546
	pricepercent	winpercent				
Reese's Peanut Butter cup	0.651	84.18029				
Reese's Miniatures	0.279	81.86626				
Twix	0.906	81.64291				
Kit Kat	0.511	76.76860				
Snickers	0.651	76.67378				

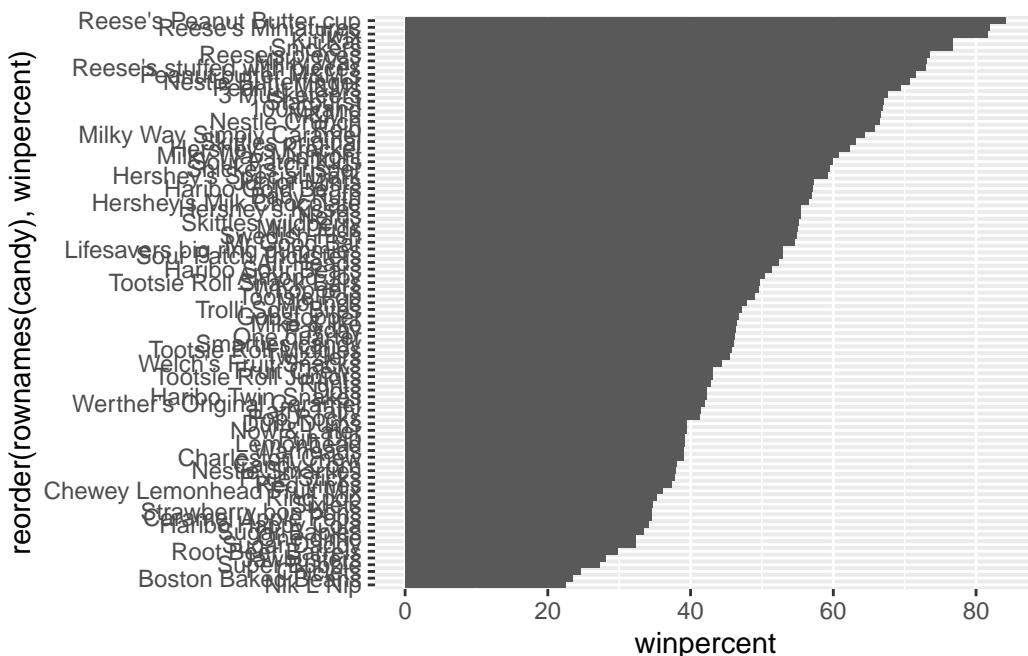
The top 5 all time favorite candy types are 1) Reese's Peanut Butter cup 2) Reese's Miniatures 3) Twix 4) Kit Kat 5) Snickers

Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy) +  
  aes(winpercent, rownames(candy), winpercent) +  
  geom_col()
```



```
ggplot(candy) +  
  aes(winpercent, reorder(rownames(candy),winpercent)) +  
  geom_col()
```

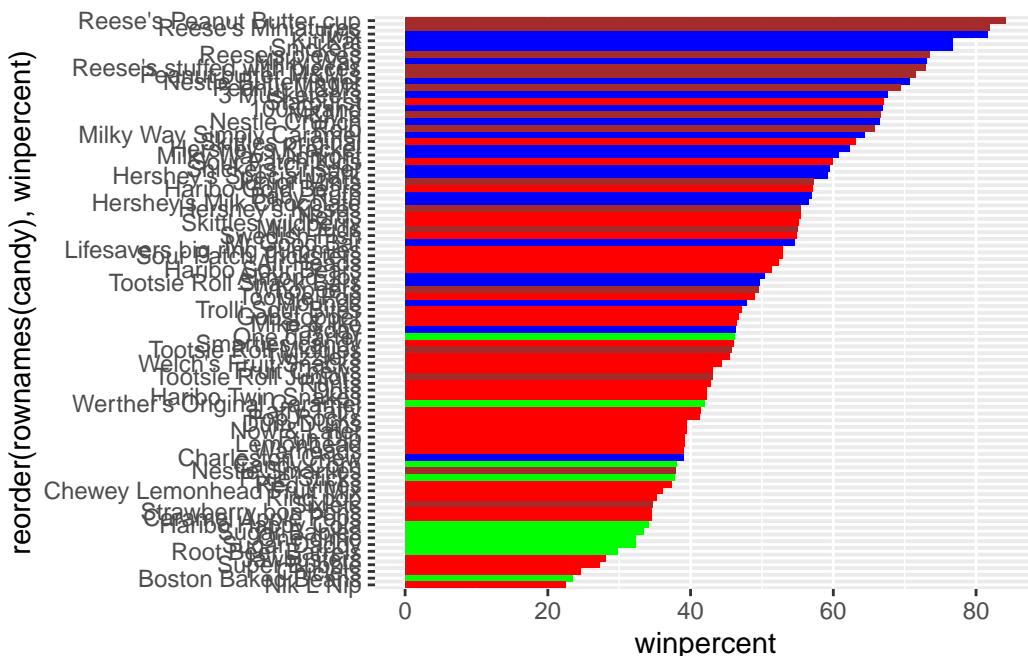


```

my_cols=rep("green", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "brown"
my_cols[as.logical(candy$bar)] = "blue"
my_cols[as.logical(candy$fruity)] = "red"

ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols)

```



Q17. What is the worst ranked chocolate candy?

Sixlets

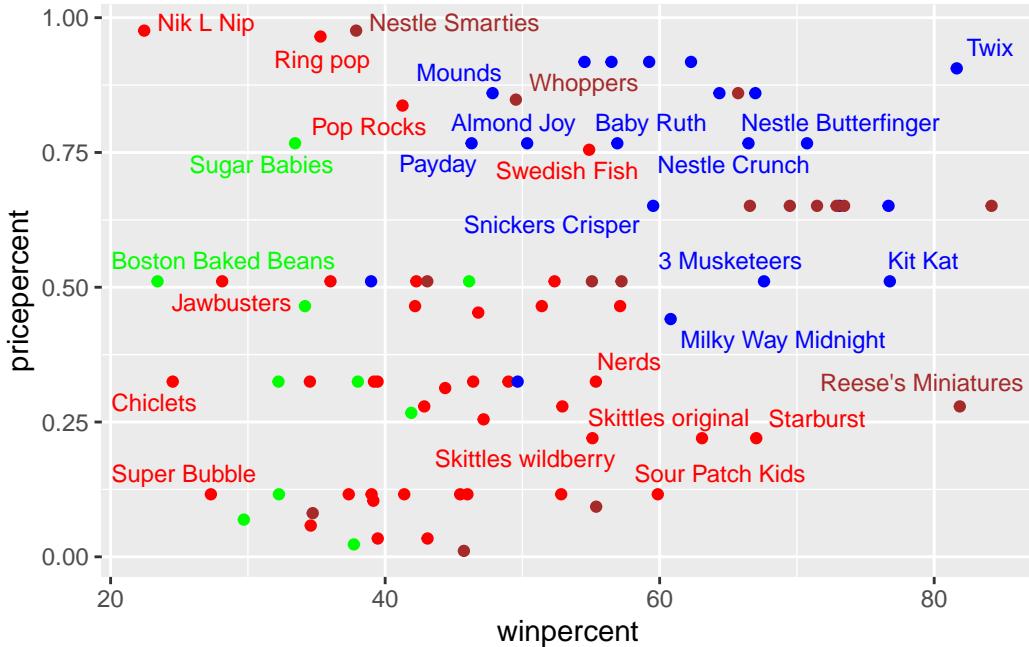
Q18. What is the best ranked fruity candy?

Starbursts

```
library(ggrepel)

ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 7)
```

Warning: ggrepel: 57 unlabeled data points (too many overlaps). Consider increasing max.overlaps



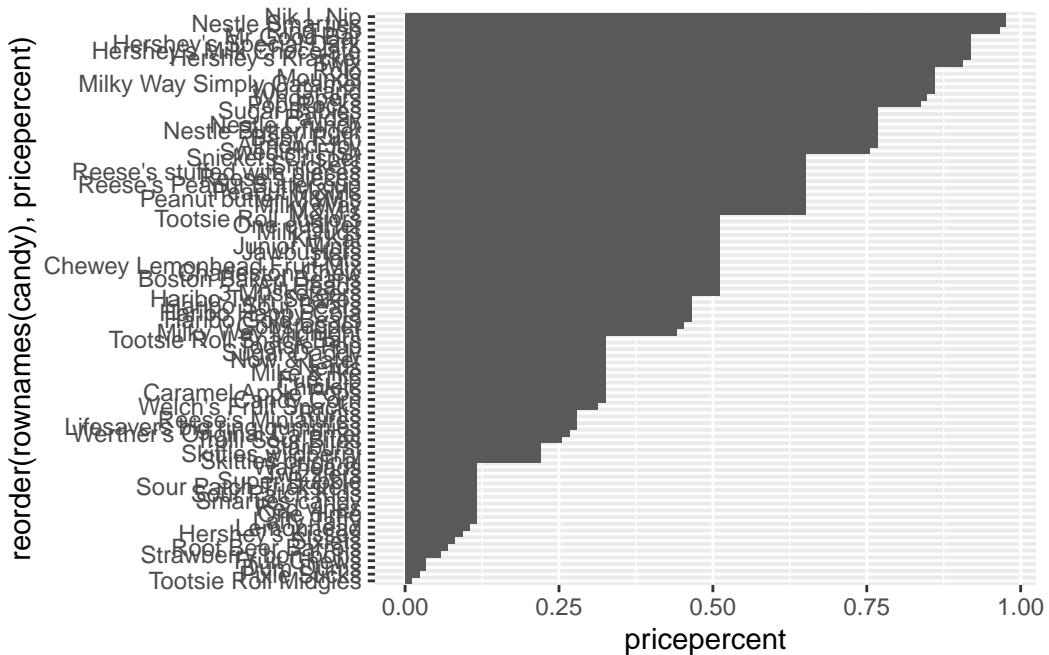
Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reese's Miniatures

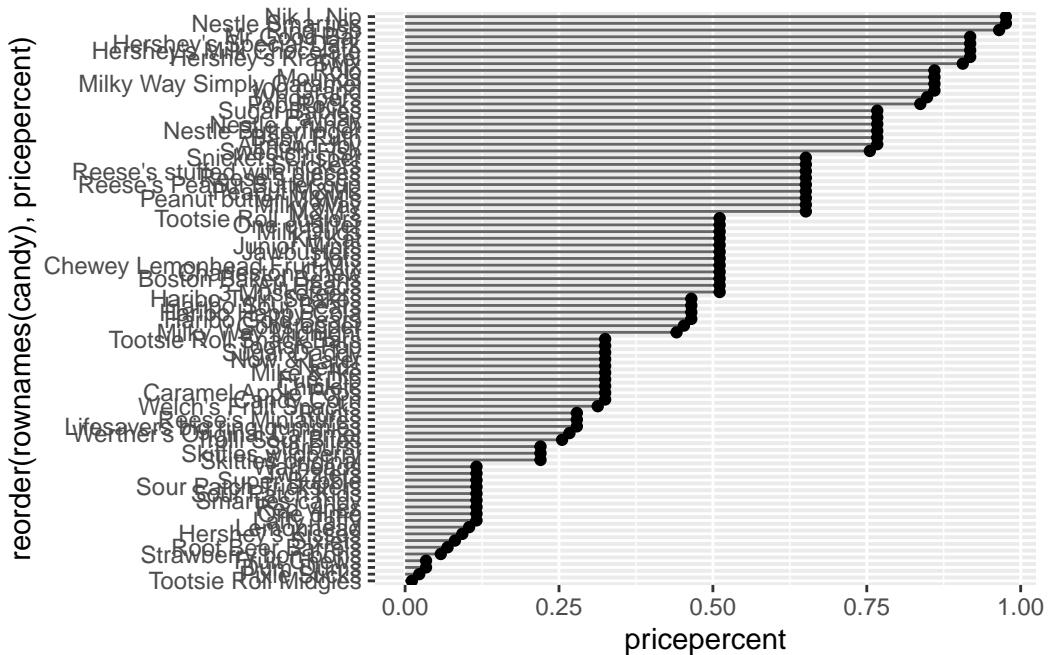
Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

Nik L Nip

```
ggplot(candy) +
  aes(pricepercent, reorder(rownames(candy), pricepercent)) +
  geom_col()
```



```
ggplot(candy) +
  aes(pricepercent, reorder(rownames(candy), pricepercent)) +
  geom_segment(aes(yend = reorder(rownames(candy), pricepercent),
                  xend = 0), col="gray40") +
  geom_point()
```

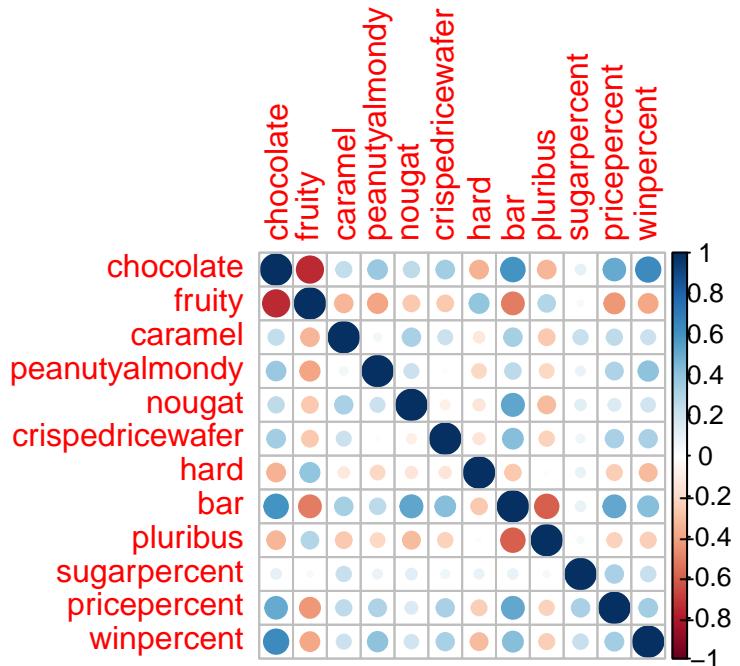


6 Exploring the coorelation structure

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)
corrplot(cij)
```



Principal Component Analysis

Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and Fruity

Q23. Similarly, what two variables are most positively correlated?

Chocolate and win percent

```
pca <- prcomp(candy, scale=TRUE)

summary(pca)
```

Importance of components:

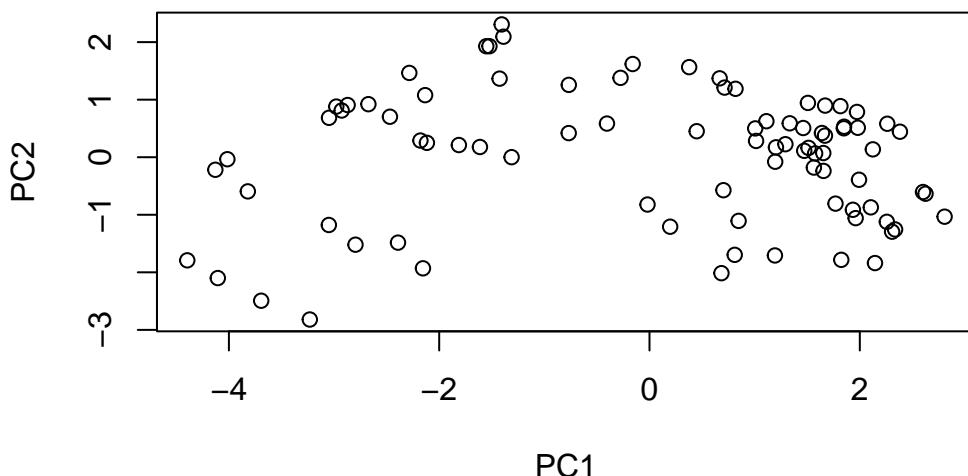
	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		

```
Standard deviation      0.74530 0.67824 0.62349 0.43974 0.39760
Proportion of Variance 0.04629 0.03833 0.03239 0.01611 0.01317
Cumulative Proportion  0.89998 0.93832 0.97071 0.98683 1.00000
```

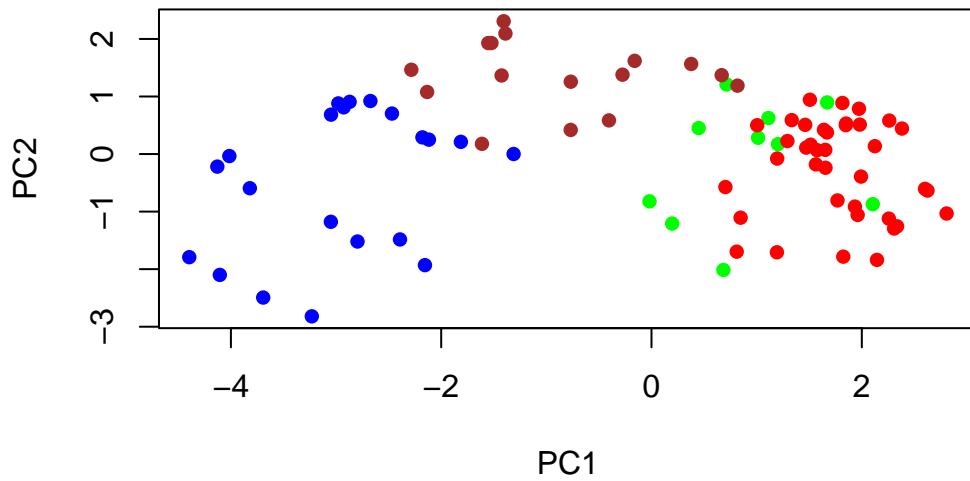
```
pca$rotation[,1]
```

chocolate	fruity	caramel	peanutyalmond
-0.4019466	0.3683883	-0.2299709	-0.2407155
nougat	crispedrice	wafer	hard
-0.2268102	-0.2215182	0.2111587	bar
pluribus	sugarpercent	pricepercent	winpercent
0.2600041	-0.1083088	-0.3207361	-0.3298035

```
plot(pca$x[,1:2])
```



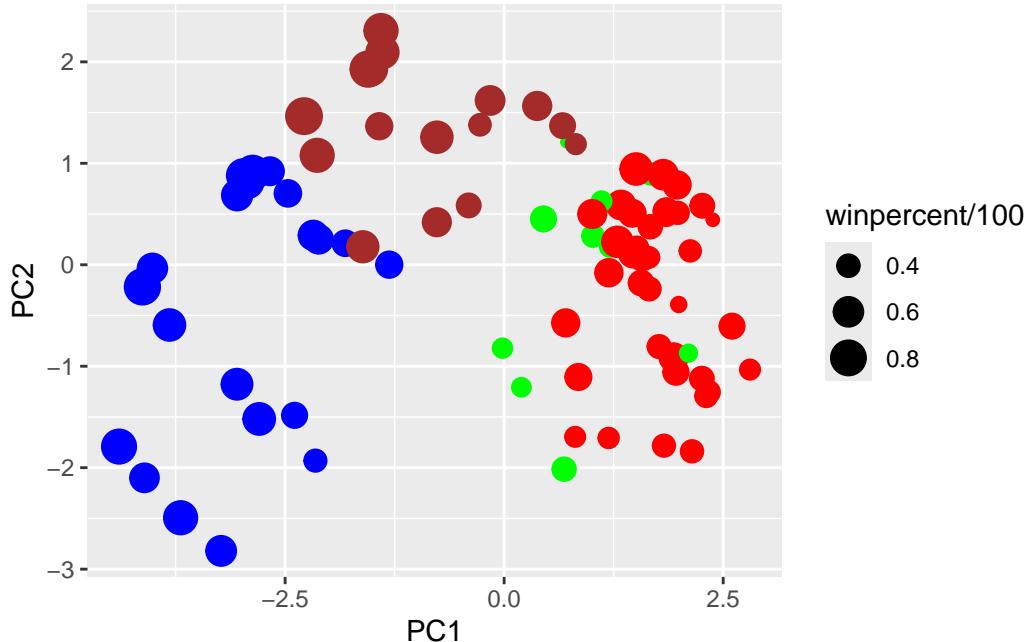
```
plot(pca$x[,1:2], col=my_cols, pch=16)
```



```
my_data <- cbind(candy, pca$x[,1:3])

p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)

p
```



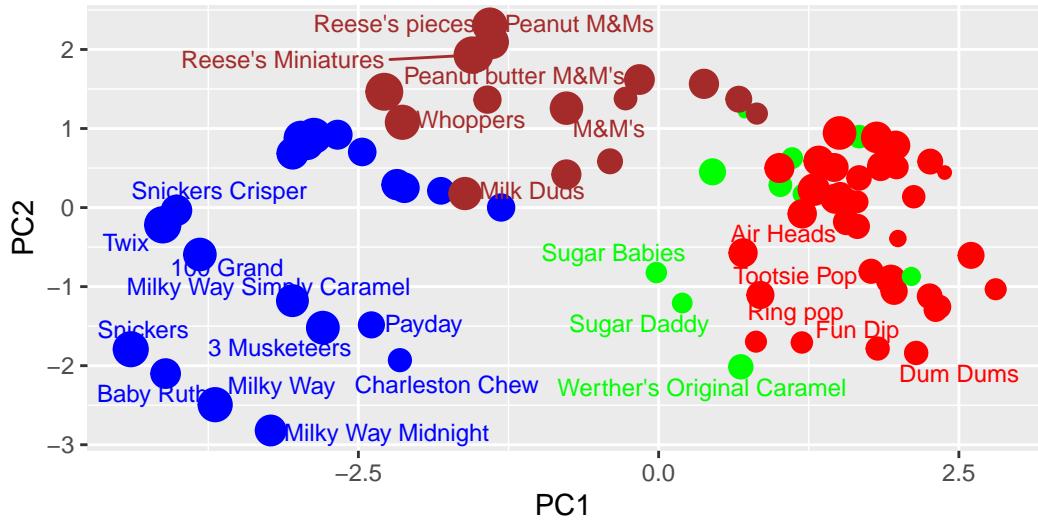
```
library(ggrepel)

p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),
       caption="Data from 538")
```

Warning: ggrepel: 59 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Halloween Candy PCA Space

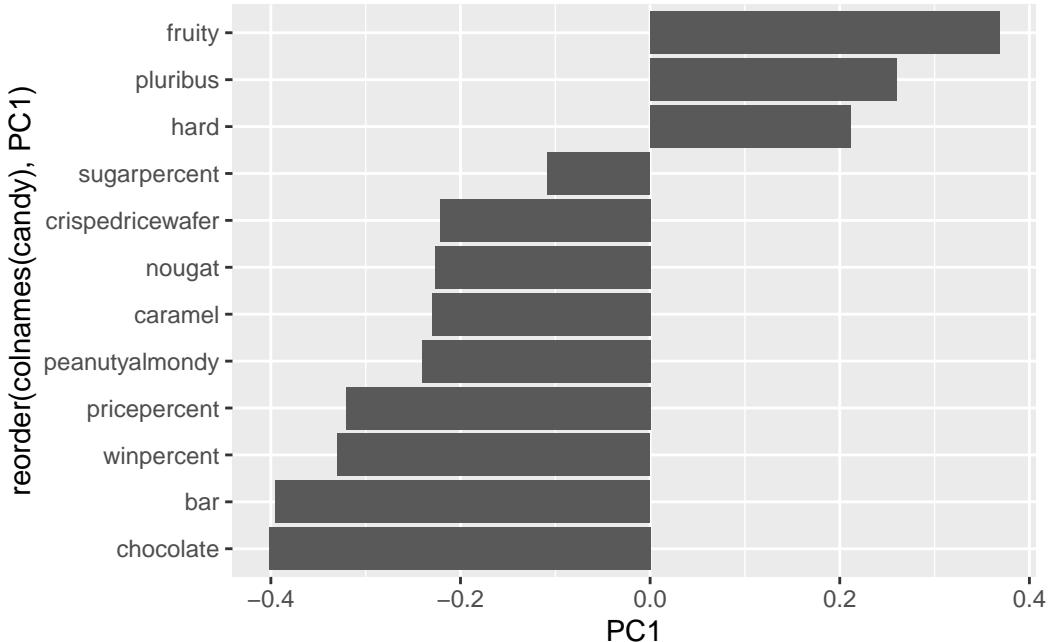
Colored by type: chocolate bar (dark brown), chocolate other (light brown),



Data from 538

Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

```
ggplot(pca$rotation) +
  aes(PC1, reorder(colnames(candy), PC1)) +
  geom_col()
```



In the positive direction fruity, Sugar percent, Crispedricewafer are seen in the positive direction this makes sense as we saw this previously in the scatter plot. This makes sense to be as we can see chocolate, bar and other variables that were negatively coorelated to fruity in the negative direction.

Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

The combination of characteristics that make something appear to be a “winning candy” would be that they are chocolate, they are a bar, and have a high price percent. The bar plot of candy ranking based on win percent values showed that bar and chocolate candy were clearly stronger than Fruity and others. price percent vs Win percent graph showed that Bar and chocolate items were both winners and more highly priced.the PCA graphs showed us the positive correlation between Chocolate, Bar, Win percent and pricepercent, which were negatively correlated with fruity, pluribus and hard.

Optional extension questions

Q26. Are popular candies more expensive? In other words: is price significantly different between “winners” and “losers”? List both average values and a P-value along with your answer.

```
losers = candy[which(candy$winpercent < 50),]  
  
winners = candy[which(candy$winpercent >= 50),]  
  
#losers %>% arrange(pricepercent)  
  
#winners %>% arrange(pricepercent)  
  
t.test(winners,losers)
```

Welch Two Sample t-test

```
data: winners and losers  
t = 2.267, df = 751.43, p-value = 0.02367  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 0.2860267 3.9814219  
sample estimates:  
mean of x mean of y  
5.619883 3.486159
```

Yes, price significantly effects the difference between winners and losers.